



MALIGNANT COMMENT CLASSIFICATION

Submitted by:
Asheem Siwach

ACKNOWLEDGMENT

I would like to thank Flip Robo Technologies for providing me with the opportunity to work on this project from which I have learned a lot. I am also grateful to Mr. Shubham Yadav for his constant guidance and support.

Reference sources are :-

- Google
- Stackoverflow.com
- Analytics vidhya
- Notes and repository from DataTrained

Research papers –

<https://www.nltk.org/book/ch05.html>

<https://www.analyticsvidhya.com/blog/2017/01/ultimate-guide-to-understand-implement-natural-language-processing-codes-in-python/>

INTRODUCTION

- Business Problem Framing

The proliferation of social media enables people to express their opinions widely online. However, at the same time, this has resulted in the emergence of conflict and hate, making online environments uninviting for users. Although researchers have found that hate is a problem across multiple platforms, there is a lack of models for online hate detection.

Online hate, described as abusive language, aggression, cyberbullying, hatefulness and many others has been identified as a major threat on online social media platforms. Social media platforms are the most prominent grounds for such toxic behaviour.

There has been a remarkable increase in the cases of cyberbullying and trolls on various social media platforms. Many celebrities and influences are facing backlashes from people and have to come across hateful and offensive comments. This can take a toll on anyone and affect them mentally leading to depression, mental illness, self-hatred and suicidal thoughts.

- Conceptual Background of the Domain Problem

Internet comments are bastions of hatred and vitriol. While online anonymity has provided a new outlet for aggression and hate speech, machine learning can be used to fight it. The problem we sought to solve was the tagging of internet comments that are aggressive towards other users. This means that insults to third parties such as celebrities will be tagged as unoffensive, but “u are an idiot” is clearly offensive.

Our goal is to build a prototype of online hate and abuse comment classifier which can used to classify hate and offensive comments so that it can be controlled and restricted from spreading hatred and cyberbullying.

Analytical Problem Framing

- Mathematical/ Analytical Modeling of the Problem

This was a NLP Project and in this project we deal with the textual data and for understanding the data we used some methods like removing punctuations, numbers, stopwords and using the lemmatization process to convert the complex words into their simpler forms. These processes helped in cleaning the unwanted words from the comments and we were left with only those words which were going to help in our model building. After cleaning the data we used TF-IDF Vectorizer technique to convert textual data into vector form. This technique works on the basis of the frequency of words present in the document. After training with the train dataset we use this technique on the test dataset.

- Data Sources and their formats

This data was provided to me by FlipRobo Technologies in a csv file format. This file contains training and testing datasets. On the training dataset we build a model and using this model we have to predict the outcomes from the testing dataset.

The data set contains the training set, which has approximately 1,59,000 samples and the test set which contains nearly 1,53,000 samples. All the data samples contain 8 fields which include 'Id', 'Comments', 'Malignant', 'Highly malignant', 'Rude', 'Threat', 'Abuse' and 'Loathe'.

The label can be either 0 or 1, where 0 denotes a NO while 1 denotes a YES. There are various comments which have multiple labels. The first attribute is a unique ID associated with each comment.

The data set includes:

-**Malignant:** It is the Label column, which includes values 0 and 1, denoting if the comment is malignant or not.

-**Highly Malignant:** It denotes comments that are highly malignant and hurtful.

-**Rude:** It denotes comments that are very rude and offensive.

-**Threat:** It contains indication of the comments that are giving any threat to someone.

-**Abuse:** It is for comments that are abusive in nature.

-**Loathe:** It describes the comments which are hateful and loathing in nature.

-ID: It includes unique Ids associated with each comment text given.

-Comment text: This column contains the comments extracted from various social media platforms.

First of all we upload the training dataset into df_train dataframe and check the datatypes present in it.

```
In [4]: features_info(df_train)

Rows in dataset = 159571

Columns in dataset = 8

Features names =
Index(['id', 'comment_text', 'malignant', 'highly_malignant', 'rude', 'threat',
      'abuse', 'loathe'],
      dtype='object')

Dataset types :
id                object
comment_text      object
malignant         int64
highly_malignant  int64
rude              int64
threat            int64
abuse             int64
loathe            int64
dtype: object
```

So the data present in training data is both object and integer in nature. 2 columns are object dtypes and 6 columns are integer in dtypes.

• Data Preprocessing Done

Data preprocessing is the data mining technique that involves transforming raw data into an understandable data format. So in data preprocessing technique first step is import data and the libraries to be used in model building.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re

from nltk import pos_tag
from nltk.stem import WordNetLemmatizer
from gensim.parsing.preprocessing import STOPWORDS
from nltk.corpus import wordnet
import warnings
warnings.filterwarnings('ignore')

# Importing libraries for classification
from sklearn.model_selection import train_test_split, cross_val_score

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
```

Loading Dataset

```
In [2]: # Loading training dataset
df_train = pd.read_csv(r"D:\Flip robo inter\data\Malignant-Comments-Classifer-Project--1\Malignant Co
df_train.head(5)
```

```
Out[2]:
```


	id	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe
0	0000997932d777bf	Explanation\nWhy the edits made under my usern...	0	0	0	0	0	0
1	000103f0d9cfb60f	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0
3	0001b41b1c6bb37e	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	0
4	0001d958c54c6e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0

Now after loading the dataset we check the missing values –

```
In [3]: # Creating function to see information of dataset
def features_info(dataset):

    print("\nRows in dataset =",dataset.shape[0])
    print("\nColumns in dataset =",dataset.shape[1])
    print("\nFeatures names =",dataset.columns)
    print("\nDataset types : \n",dataset.dtypes)
    print("\nNull values in dataset : \n",dataset.isnull().sum())
    print("\nAny duplicated values :",dataset.duplicated().values.any(),"\n")
    for i in dataset.columns:
        print("Total unique values in {} = {}".format(i,dataset[i].nunique()))
    plt.figure(figsize=(10,7))
    sns.heatmap(dataset.isnull())
    plt.title("Graph showing null values",color='steelblue',fontsize=15)
    plt.show()
```

```
In [4]: features_info(df_train)
```



```
In [17]: # Creating function for cleaning data

def cleaning_data(text):
    # Replace email addresses with 'email'
    text=re.sub(r'^.+@[^\.\.]*\.[a-z]{2,}$','email', text)

    # Replace 10 digit phone numbers (formats include paranthesis, spaces, no spaces, dashes) with 'phonenummer'
    text=re.sub(r'^\((?\d{3})?\s-?\d{3}\s-?\d{4}$','phonenummer',text)

    # converting text into lower case
    text = text.lower()

    # Removing all the special characters
    text=re.sub(r'^\w+', " ",text)

    # getting only words(i.e removing all the" _ ")
    text=re.sub(r'[\_]+', " ",text)

    # getting rid of unwanted characters(i.e remove all the single characters left)
    text=re.sub(r'\s+[a-zA-Z0-9]\s+', " ",text)

    # Remove number
    text=re.sub(r'\d+(\.\d+)?','',text)

    # Removing extra whitespaces
    text = re.sub(r'\s+', " ",text ,flags=re.I)

    return text

In [18]: # text column after removing unwanted characters
df_train['clean_text1'] = df_train['comment_text'].apply(cleaning_data)
```

After the removal of unwanted notations we moved to remove **stopwords** from our dataset. For removing them we created another function named `stop_words` and append the text into another new column.

```
In [19]: # Lets remove stopwords using gensim library
def stop_words(text):
    # splitting the text
    text= text.split()

    #stopwords removal
    text = [w for w in text if w not in set(STOPWORDS)]

    return text

In [20]: # text column after removing stopwords
df_train['clean_text2'] =df_train['clean_text1'].apply(stop_words)
```

After removing all the unnecessary words or numbers we converted the words into their simpler form using the **Lemmatization process**. In this process we defined two function , first one will tag the words into proper format and the other function will convert them into simpler form using the tagged alphabet.

```

In [22]: # creating function to rename pos_tags
def get_wordnet_pos(text_tag):
    if text_tag.startswith('J'):
        return wordnet.ADJ
    elif text_tag.startswith('V'):
        return wordnet.VERB
    elif text_tag.startswith('N'):
        return wordnet.NOUN
    elif text_tag.startswith('R'):
        return wordnet.ADV
    else:
        return wordnet.NOUN

In [23]: # Now creating function for the Lemmatization
def lemmatization(text):
    # getting pos tags
    text_tags = pos_tag(text)

    # Now do Lemmatization
    text = [(WordNetLemmatizer().lemmatize(w[0], get_wordnet_pos(w[1])) for w in text_tags)]

    # joining the tokens
    text = ' '.join(text)

    return text

In [24]: # fully cleaned textual column
df_train['cleaned_comment'] = df_train['clean_text2'].apply(lemmatization)

```

Now we took a random sample from our dataset and compared the text of before and after the treatment.

```

In [26]: # Sample before cleaning
df_train['comment_text'][3]

Out[26]: '\nMore\nI can\'t make any real suggestions on improvement - I wondered if the section statistics should be later on, or a sub
section of ""types of accidents"" -I think the references may need tidying so that they are all in the exact same format ie da
te format etc. I can do that later on, if no-one else does first - if you have any preferences for formatting style on referenc
es or want to do it yourself please let me know.\n\nThere appears to be a backlog on articles for review so I guess there may b
e a delay until a reviewer turns up. It\'s listed in the relevant form eg Wikipedia:Good_article_nominations#Transport "'

In [27]: # sample after cleaning
df_train['cleaned_comment'][3]

Out[27]: 'real suggestion improvement wonder section statistic later subsection type accident think reference need tidy exact format dat
e format later preferences format style reference want let know appear backlog article review guess delay reviewer turn list re
levant form wikipedia good article nomination transport'

```

Encoding the categorical data (Feature Extraction and scaling)

As of now we had cleaned the data , now another major step towards building a model is to convert the textual data into numerical form because our algorithms understand only the numerical data. So for converting into numerical or vector form we used Tf-Idf Vectorizer technique which converted the textual data into the vectors using the terms frequency method.


```

In [37]: # Convert text data into vector form
from sklearn.feature_extraction.text import TfidfVectorizer

tf_idf = TfidfVectorizer(min_df=4)
# fitting
fitting_idf = tf_idf.fit(df_train1['cleaned_comment'])

In [38]: #Due to insufficient memory we can not use dataframe method to see all the features extracted.

def vectorization(fitting_data):

    return tf_idf.transform(df_train1['cleaned_comment'])

In [39]: # independent features
x=vectorization(fitting_idf)
x.shape

Out[39]: (159362, 34312)

In [40]: # Dependent feature
y = df_train1['negative_cmnts']
y.shape

Out[40]: (159362,)

```

Splitting the dataset –

The last step for data-preprocessing is splitting the dataset into training and testing dataset Using the train_test_split model selection method we converted the dataset into training and testing.

Splitting dataset into training and testing

```

In [41]: # Let's split the train and test dataset
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=45,stratify=y)
print("Train shapes : X = {}, y = {}".format(x_train.shape,y_train.shape))
print("Test shapes : X = {}, y = {}".format(x_test.shape,y_test.shape))
# Success
print("\nTraining and testing split was successful.")

Train shapes : X = (111553, 34312), y = (111553,)
Test shapes : X = (47809, 34312), y = (47809,)

Training and testing split was successful.

```

• Data Inputs- Logic- Output Relationships

```

# Let's plot the correlation chart

df_train.corr().style.background_gradient(cmap='rocket')

```

	malignant	highly_malignant	rude	threat	abuse	loathe
malignant	1.000000	0.308619	0.676515	0.157058	0.647518	0.266009
highly_malignant	0.308619	1.000000	0.403014	0.123601	0.375807	0.201600
rude	0.676515	0.403014	1.000000	0.141179	0.741272	0.286867
threat	0.157058	0.123601	0.141179	1.000000	0.150022	0.115128
abuse	0.647518	0.375807	0.741272	0.150022	1.000000	0.337736
loathe	0.266009	0.201600	0.286867	0.115128	0.337736	1.000000

All the harmful comments present in dataset are having positive correlation with each others. Rude and abuse are highly correlated.

- Hardware and Software Requirements and Tools Used

There is no such requirement for hardware ,but I have used intel i5 8th generation processor with 6 GB Ram.

Software: Jupyter Notebook (Anaconda 3)

Language : Python 3.9

Libraries used in project:

- a. Pandas
- b. Numpy
- c. Matplotlib
- d. Seaborn
- e. Sklearn

Model/s Development and Evaluation

- Identification of possible problem-solving approaches (methods)

In this project there were 5-6 features which defines the type of comment like malignant , hate , abuse, threat, loathe but we created another feature named as “negative_cmnts” which is combined of all the above features and contains the labelled data into the format of 0 and 1 where 0 represents “NO” and 1 represents “Yes”.

As we have labelled data into our target feature which is the case of classification method. So we are going to use algorithms based on classification.

- Testing of Identified Approaches (Algorithms)

Based on the classification approach we are going to use following approaches :

- I. Logistic Regression
- II. Decision Tree Classifier
- III. RandomForest Classifier
- IV. AdaBoost Classifier
- V. GradientBoosting Classifier

In NLP we use Naïve Bayes algorithms mostly but due to our systems we faced memory error to deal with them.

- Run and Evaluate selected models

1. Logistic Regression -

[illegible]

II. DecisonTree Classifier –

```
classification(DecisionTreeClassifier())
```

```
DecisionTreeClassifier()
```

Training Score 99.94 %

Accuracy score = 94.17 %
Log loss = 2.01

```
Confusion matrix :
[[41613  1329]
 [ 1458 3409]]
```

Classification report :					
	precision	recall	f1-score	support	
	0	0.97	0.97	0.97	42942
	1	0.72	0.70	0.71	4867
	accuracy			0.94	47809
	macro avg	0.84	0.83	0.84	47809
	weighted avg	0.94	0.94	0.94	47809

Cross val score = 83.913 %
Standard deviation = 0.0020

ROC AUC : 83.4741377882354

[illegible]

III. RandomForest Classifier –

```
classification(RandomForestClassifier())
```

```
RandomForestClassifier()
```

Training Score 99.94 %

Accuracy score = 95.65 %
Log loss = 1.50

```
Confusion matrix :
[[42588  354]
 [ 1728 3139]]
```

Classification report :				
	precision	recall	f1-score	support
0	0.96	0.99	0.98	42942
1	0.90	0.64	0.75	4867
accuracy			0.96	47809
macro avg	0.93	0.82	0.86	47809
weighted avg	0.95	0.96	0.95	47809

```
Cross val score = 96.47 %
Standard deviation = 0.0013
```

ROC AUC : 81.83560737124918

IV. AdaBoost Classifier –

```
classification(AdaBoostClassifier())
```

```
AdaBoostClassifier()
```

Training Score 94.78 %

Accuracy score = 94.62 %
Log loss = 1.86

```
Confusion matrix :
[[42507  435]
 [ 2138 2729]]
```

Classification report :		recall	f1-score	support
	precision			
0	0.95	0.99	0.97	42942
1	0.86	0.56	0.68	4867
accuracy			0.95	47809
macro avg			0.83	47809
weighted avg			0.94	47809

```
Cross val score = 89.769 %
Standard deviation = 0.0031
```

ROC AUC : 77.52925384028919

V. GradientBoosting Classifier –

```
classification(GradientBoostingClassifier())
```

```
GradientBoostingClassifier()
```

Training Score 94.37 %

Accuracy score = 94.15 %
Log loss = 2.02

```
Confusion matrix :
[[42815  127]
 [ 2670 2197]]
```

Classification report :		recall	f1-score	support
	precision			
0	0.94	1.00	0.97	42942
1	0.95	0.45	0.61	4867
accuracy			0.94	47809
macro avg	0.94	0.72	0.79	47809
weighted avg	0.94	0.94	0.93	47809

```
Cross val score = 90.066 %
Standard deviation = 0.0026
```

ROC AUC : 72.42249801594473

[illegible]

Logistic Regression Hypertunning

```
In [55]: # Defining parameters
solvers = ['newton-cg', 'lbfgs', 'liblinear']
penalty = ['l2']
c_values = [100, 10, 1.0, 0.1, 0.01]

In [58]: from sklearn.model_selection import RepeatedStratifiedKFold

# define grid search
grid = dict(solver=solvers, penalty=penalty, C=c_values)
# cross validation technique
cv_method = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)

lg_hyper = RandomizedSearchCV(estimator=LogisticRegression(), param_distributions=grid, n_jobs=-1, cv=cv_method)
lg_hyper_result = lg_hyper.fit(x, y)

Best: 0.960419 using {'solver': 'lbfgs', 'penalty': 'l2', 'C': 10}
0.956878 (0.001116) with: {'solver': 'liblinear', 'penalty': 'l2', 'C': 1.0}
0.960419 (0.001314) with: {'solver': 'lbfgs', 'penalty': 'l2', 'C': 10}
0.956880 (0.001124) with: {'solver': 'lbfgs', 'penalty': 'l2', 'C': 1.0}
0.954165 (0.001351) with: {'solver': 'newton-cg', 'penalty': 'l2', 'C': 100}
0.960386 (0.001338) with: {'solver': 'newton-cg', 'penalty': 'l2', 'C': 10}
0.954176 (0.001326) with: {'solver': 'liblinear', 'penalty': 'l2', 'C': 100}
0.907270 (0.000863) with: {'solver': 'lbfgs', 'penalty': 'l2', 'C': 0.01}
0.937815 (0.001106) with: {'solver': 'newton-cg', 'penalty': 'l2', 'C': 0.1}
0.907368 (0.000860) with: {'solver': 'liblinear', 'penalty': 'l2', 'C': 0.01}
0.937812 (0.001103) with: {'solver': 'lbfgs', 'penalty': 'l2', 'C': 0.1}
```

Accuracy score = 96.0 %


Cross validation score: 95.96829843280807

Confusion matrix

```
[[42479  463]
 [ 1452 3415]]
```

Classification report

	precision	recall	f1-score	support
0	0.97	0.99	0.98	42942
1	0.88	0.70	0.78	4867
accuracy			0.96	47809
macro avg	0.92	0.85	0.88	47809
weighted avg	0.96	0.96	0.96	47809



• Key Metrics for success in solving problem under consideration

For solving the problems and understanding the result of each algorithms we have used a lot of metrics :



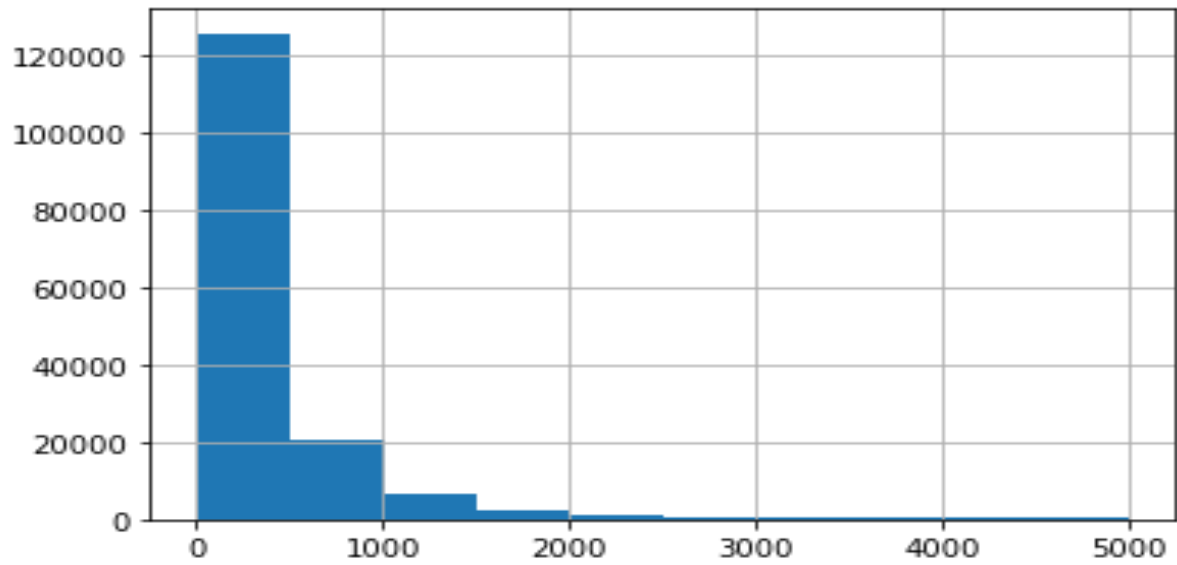
- Accuracy Score
- Classification Report
- Confusion Matrix
- Log Loss
- Roc-Auc




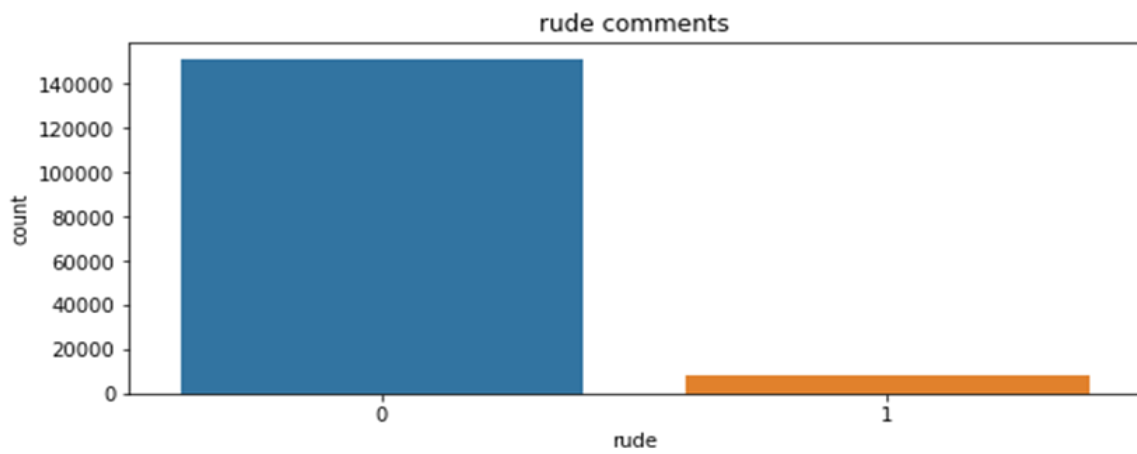
• Visualizations



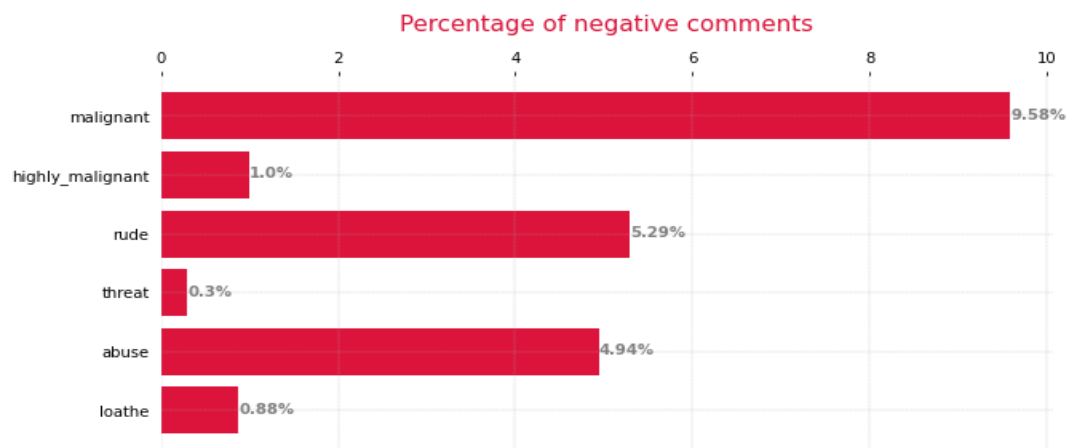
For the Visualization we have used Matplotlib and Seaborn library to plot the numerical data into graphs –

```
In [8]: # Distribution plot of comment length before cleaning  
cmnt_len_train.hist()
```

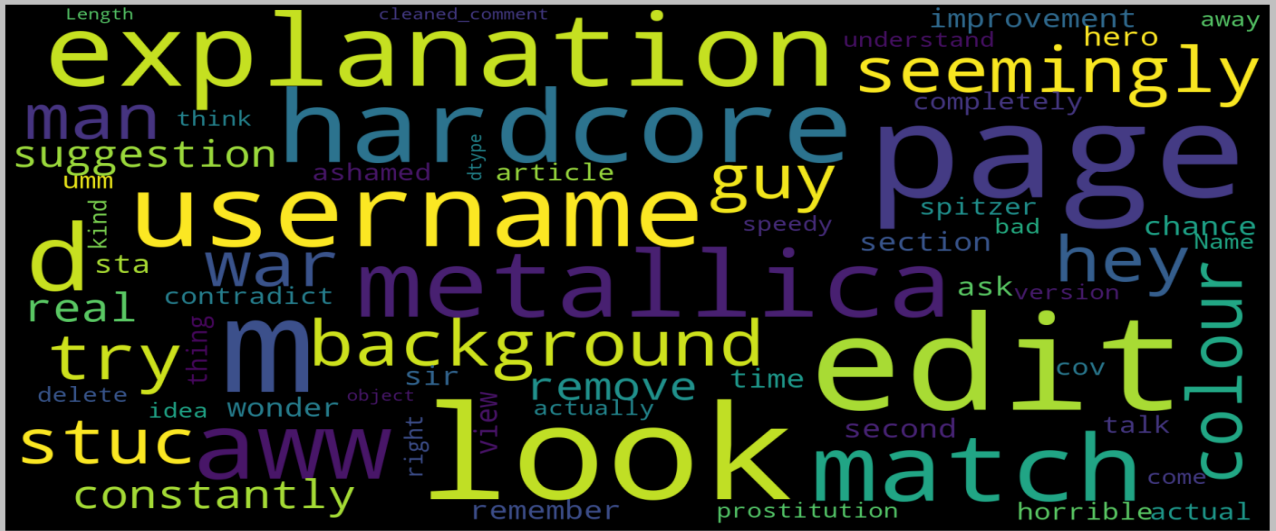
```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x23115cb82b0>
```








Words Present in Postive Comments



Words present in Negative Comments



Words Present in Testing dataset comments



• Interpretation of the Results

In [49]:

```
# Table view of result of each metrix from above algorithms
evaluations = pd.DataFrame({"Model":Model,"Accuracy":Accuracy,"Log_loss":Log_loss,"ROC_AUC":ROC_AUC,
                           "CV Score":CVScore,"Std_dev":Std})
evaluations
```

Out[49]:

	Model	Accuracy	Log_loss	ROC_AUC	CV Score	Std_dev
0	LogisticRegression()	95.55	1.54	0.800616	97.189	0.000967
1	DecisionTreeClassifier()	94.17	2.01	0.834741	83.913	0.001973
2	RandomForestClassifier()	95.65	1.50	0.818356	96.470	0.001312
3	AdaBoostClassifier()	94.62	1.86	0.775293	89.769	0.003111
4	GradientBoostingClassifier()	94.15	2.02	0.724225	90.066	0.002586

As our target feature is imbalanced which means that alone accuracy score will not give best results but instead accuracy we are using the classification report log loss score , Roc-Auc score and confusion matrix to find the best algorithm which is above all.

So we selected Logistic Regression and RandomForest Classifier based on above condition and use RandomizedSearchCV for hyper tuning the parameters of these 2 selected algorithms . After hyper tuning the parameters we selected the Logistic Regression algorithm as it give upto 85% score of Roc-Auc and losgg loss of 1.38 far better than RandomForest Classifier whose log loss was above 2.0 and again the classification report also tells that it is better than RandomForest Classifier.

Finalising the Model

- So after tuning the parameters of both the selected algorithms we are finalised Logistic Regression for our model. As the ROC AUC curve and the score of Logg loss is far better than Random Forest Classifier.
- Logistic Regression gives us Log loss of 1.38 and the Area under the curve is 85% with an Accuracy of 96%.

• Predictions

```
In [88]: df_test[['comment_text', 'Predictions']]
```

```
Out[88]:
```

	comment_text	Predictions
0	Yo bitch Ja Rule is more succesful then you'll...	1
1	== From RfC == \n\n The title is fine as it is...	0
2	" \n\n == Sources == \n\n * Zawe Ashton on Lap...	0
3	:If you have a look back at the source, the in...	0
4	I don't anonymously edit articles at all.	0
...
153159	. \n i totally agree, this stuff is nothing bu...	1
153160	== Throw from out field to home plate. == \n\n...	0
153161	" \n\n == Okinotorishima categories == \n\n I ...	0
153162	" \n\n == ""One of the founding nations of the...	0
153163	" \n :::Stop already. Your bullshit is not wel...	1

153164 rows × 2 columns

These are the predictions rom our testing dataset.

```
In [86]: # Let's see the value counts
df_test['Predictions'].value_counts()
```

```
Out[86]: 0      121873
         1       31291
         Name: Predictions, dtype: int64
```

CONCLUSION

• Key Findings and Conclusions of the Study

The conclusion for our study :-

- In training dataset we have only 10% of data which is spreading hate on social media.
- In this 10% data most of the comments are malignant ,rude or abuse.
- After using the wordcloud we find that there are so many abusive words present in the negative comments. While in positive comments there is no use of such comments.
- Some of the comments are very long while some are very short.

- Learning Outcomes of the Study in respect of Data Science

From this project we learned a lot. Gains new techniques and ways to deal with uncleaned data. Find a solution to deal with multiple target features. Tools used for visualizations gives a better understanding of dataset. We have used a lot of algorithms and find that in the classification problem where we have only two labels , Logistic Regression gives better results compared to others.

But due to our system we could not use algorithms which gives much better results in NLP project like GaussinaNB, MultinomailNM. We also used googlecolab and some pipelines techniques but none of them worked here and also it was too much time consuming.

- Limitations of this work and Scope for Future Work

This project was amazing to work on , it creates new ideas to think about but there were some limitations in this project like unbalanced dataset, multiple target features. To overcome these limitations we have to use balanced dataset so that during the training of dataset our algorithm will not give biased result.

Thankyou